Citation Integrity Audit

AI-Assisted Research Paper — Lab 1

Student: Ujjwal Gupta
Roll Number: MDE2026004
Programme: M.Tech in IT (Data Engineering)
Topic: Failure modes in Chained Tool-Calling Agents Operating Without Human Checkpoints
AI Tool: ChatGPT | **Access:** Free | **Web browsing:** Yes | **Deep Research:** No

1. Purpose of the Audit

This audit evaluates citation authenticity and integrity in the AI-assisted research paper. It checks whether selected references exist, whether bibliographic metadata and identifiers are consistent, and whether the cited sources support the claims for which they were used.

2. Paper and Reference Profile

Measure	Result
Approximate paper length	Approximately 10 pages
Approximate word count	Approximately 5,000 words
Major sections	10
Total references	18
Approximate in-text citation occurrences	30+
AI warned to verify references?	No

3. Systematic Audit Sample

The bibliography contains 18 references. Ten references were selected using the systematic sampling procedure: R1=1, R2=2, R3=3, R4=7, R5=9, R6=11, R7=13, R8=16, R9=17, and R10=18.

Slot	Ref.	Plausibility	Prediction	Final
R1	1	5	Genuine	A — Verified
R2	2	5	Genuine	A — Verified
R3	3	5	Genuine	A — Verified
R4	7	5	Genuine	A — Verified
R5	9	5	Genuine	A — Verified
R6	11	5	Genuine	A — Verified
R7	13	4	Genuine	A — Verified
R8	16	5	Genuine	A — Verified
R9	17	5	Genuine	A — Verified
R10	18	5	Genuine	A — Verified

4. Verification Results

Slot	Identifier / record	Finding
R1	arXiv:2302.04761	Toolformer exists; title, authors, year and identifier match.
R2	arXiv:2210.03629 / ICLR 2023	ReAct exists; title, authors and year match.
R3	arXiv:2307.13854 / ICLR 2023	WebArena exists; metadata and identifier match.
R4	DOI 10.1145/3605764.3623985	Publication exists; metadata and DOI match.
R5	NeurIPS 2024 / DOI 10.52202/079017-2636	AgentDojo metadata and DOI match.
R6	NeurIPS 2023 / arXiv:2303.11366	Reflexion metadata and identifier match.

Slot	Identifier / record	Finding
R7	ICLR 2025 / arXiv:2410.02644	Agent Security Bench metadata and identifier match.
R8	ACL 2025 / DOI 10.18653/v1/2025.acl-long.1435	Task Shield metadata, pages and DOI match.
R9	DOI 10.3390/electronics15132829	Journal metadata and DOI match.
R10	USENIX Security '26 official record	AttriGuard metadata matches; no DOI supplied.

5. Classification and Scores

All ten audited references were classified A — Verified. No sampled reference was classified as wrong metadata, Frankenstein, fabricated, or identifier mismatch.

Classification	Count	Points
A — Verified	10	4 each
B — Wrong metadata	0	3 each
C — Frankenstein	0	1 each
D — Fabricated	0	0
E — Identifier mismatch	0	1 each

Authenticity points: 4(10) + 3(0) + 1(0) + 0(0) + 1(0) = 40.

Authenticity Score: 40 / (4 × 10) × 100 = **100/100**.

Prediction Accuracy: 10 / 10 × 100 = **100%**.

6. Claim–Citation Support

For the ten audited references, the sources were found to be genuine and the audited claims were recorded as supported by the cited publications. This conclusion applies to the audited sample and is not an exhaustive verification of all 18 references.

7. Final Interpretation

Did professional-looking citations necessarily turn out to be genuine? No. Professional formatting alone was not treated as evidence; authenticity was established through verification.

Most serious citation-integrity problem observed: No fabricated or mismatched reference was found in the audited sample. The principal risk is false confidence from convincing-looking citations before verification.

Single most important lesson: A citation should be checked both for the existence and correctness of its bibliographic record and for whether the source actually supports the cited claim.

8. Audit Limitation

The paper contains 18 references, while 10 were deeply audited using the prescribed systematic sample. Therefore, the 100/100 authenticity score and 100% prediction accuracy describe the audited sample, not all 18 references.